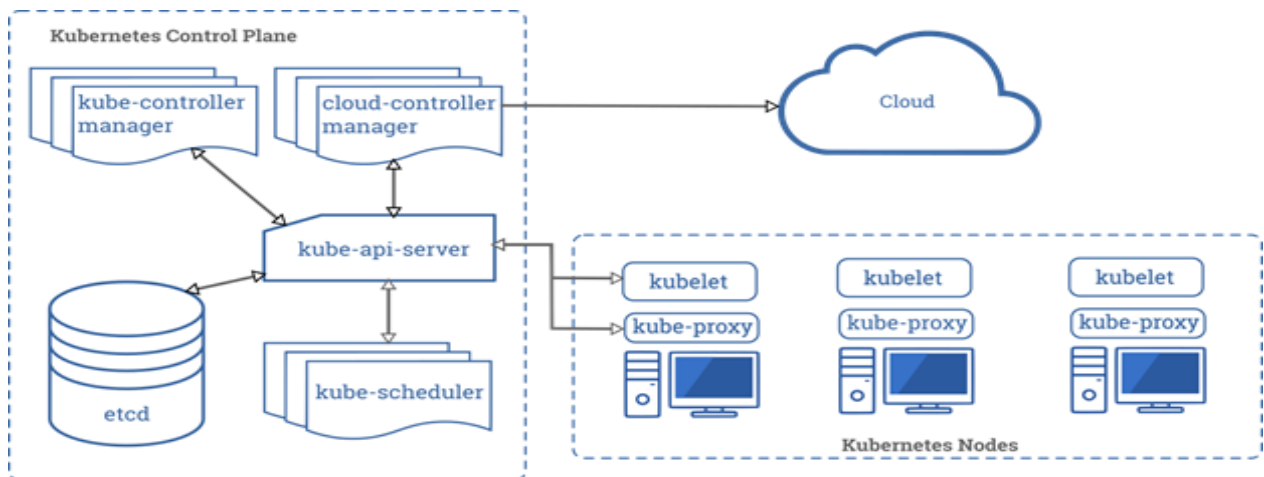


KUBERNETES

- Kubernetes is a cluster manager for docker used for orchestration, service discovery and load balancing
- It is master and node concept
- Kubernetes minions run tasks passed to it by the Kubernetes master
- Kubernetes uses Etcd to communicate master and nodes
- Kubernetes having a cluster which it manages the containers in the server when servers get down and run the containers in the servers which was up and auto scales the containers in the server which is up

Components of Kubernetes

- 1.API Server
- 2.Etcd
- 3.kubelet
- 4.scheduler
- 5.controller
- 6.Container runtime



1.API Server (application interface programming Server)

It Acts as the front-end server for kubernetes(API sever is the only way we can talk with kubernetes)

there is command called kubectl by this we can talk to API servers in Yaml

here we can create multiple api server.(odd number)

2.Etcd

Etcd is a distributed reliable key value store to store all data to manage cluster

it is rap database, all information will be saved

3.kubelet

Kubelet is an agent to runs on each node of a cluster, Responsible to make sure containers are running in the pod.

4.scheduler

Scheduler is to distribute the work across all nodes or containers and assign pods to the node

5.Kube controller manager

Main brain behind the kubernetes. If some any containers go down automatically creates new container

6.Cloud Controller Manager

It will embed cloud-specific logic, it is used when we use in cloud platform

if you're working on on-premise this service will be not used

7.Kube-Proxy

it will maintain the network rules and nodes

8.Container runtime

Container run time in our case it is docker

Different ways to Install kubernetes

1.Development environment(minikube)

2.production environment(Kubernetes Operations(KOPS))

2.production environment(Kubernetes Operations(KOPS))













First, we need to create Domain Name and link to Route 53

create s3 bucket

Create one Mgmt Server(Ec2 instance)-t2.micro is enough and enter into the server

let's go to <https://github.com/kubernetes/kops/releases>

Installation of kops

Assets 24		
	channels-linux-amd64	61.6 MB
	channels-linux-amd64.sha256	65 Bytes
	channels-linux-arm64	58.5 MB
	channels-linux-arm64.sha256	65 Bytes
	images.tar.gz	169 MB
	images.tar.gz.sha256	65 Bytes
	kops-darwin-amd64	155 MB
	kops-darwin-amd64.sha256	65 Bytes
	kops-linux-amd64 ✓	155 MB
	kops-linux-amd64.sha256	65 Bytes
	kops-linux-arm64	147 MB
	kops-linux-arm64.sha256	65 Bytes

copy the link address

```
wget <link>
```

```
chmod 700 kops-linux-amd64
```

```
mv kops-linux-amd64 kops
```

```
mv kops /user/local/bin/
```

```
kops version
```

or we can go to latest version also

Installation of Kubectl

download kubectl from browser

```
chmod 700 kubectl
```

```
mv kubectl /user/local/bin/
```

```
kubectl version
```

here KOPS is for k8s cluster Mgmt and KUBECTL for k8s Cluster ops

```
ssh-keygen
```

```
apt update && apt install unzip -y
```

AWS Cli Installation

next search in browser aws cli install ubuntu

`aws --version`

Now let's go to Aws account and go to IAM and create a user and provide programmatical access and provide admin access (or search in internet for the permissions) and copy the access key and Secret access key
now let's back to ec2 server and

`aws configure`

and enter access key and Access key and us-east-1 and json
it will show all the buckets from this we will know ur connected to aws account

Creation of cluster, Master and nodes

`kops create cluster --name=<route 53 domain> --state=s3://<bucket name> --zone=us-east-1a --node-count=2 --node-size=t2.micro --master-size=t2.small --master-volume-size 20 --node-volume-size 10 --dns-zone= <route 53 domain> --yes`

here it will create the cluster with the specific given requirements here at the end of the cmd if we didn't give any yes. it will not deploy cluster

`kops delete cluster --name=<route 53 domain> --state=s3://<bucket name> --yes`

this will delete the cluster, here we need to provide name and state

`kops get cluster --state=s3://<bucket name>`

it will show the cluster name and cloud and zones

`kops get ig --name <route 53 domain> --state s3://<bucket name>`

```
root@ip-192-168-1-210:~# kops get ig --name devopsk8s.xyz --state s3://devopsk8s.xyz
NAME                ROLE    MACHINETYPE  MIN  MAX  ZONES
master-us-east-1a   Master  t2.small     1    1    us-east-1a
nodes               Node    t2.micro      2    2    us-east-1a
```

it will show the how many masters and nodes are running and which type of instance type it is using

here ig =instance group

`kops edit ig --name<route 53 domain> <master name> --state=s3://<bucket name>`

here we can edit the master... here we need to give master name as we get in above screenshot, when we execute this cmd all the content of master will be opened and we can modify

```

apiVersion: kops.k8s.io/v1alpha2
kind: InstanceGroup
metadata:
  creationTimestamp: "2020-06-03T06:40:14Z"
  labels:
    kops.k8s.io/cluster: devopsk8s.xyz
  name: master-us-east-1a
spec:
  image: kope.io/k8s-1.17-debian-stretch-amd64-hvm-ebs-2020-01-17
  machineType: t2.small
  maxSize: 1
  minSize: 1
  nodeLabels:
    kops.k8s.io/instancegroup: master-us-east-1a
  role: Master
  rootVolumeSize: 20
  subnets:
  - us-east-1a

```

here we can edit all the data e.g.: how many master you need, how much volume you need....

`kops edit ig --name<route 53 domain> <node name> --state=s3://<bucket name>`

here we can edit nodes

`kops update cluster --name <route 53 domain> --yes --state s3://<s3 bucket name>`

after editing we need to update the cluster to make a changes deployed

`vi .bashrc`

`export NAME=devopsk8s.xyz`

`export KOPS_STATE_STORE=s3://<s3 bucket name>`

`alias <any name>='kubectl'`

`source .bash`

here we need to go to bottom of page add this steps, because if we add these here it will take kops name and state will be default, no need to specify every time in the cmd like before and enter source .bash to execute this

`kops get cluster`

from here no need of providing name and state because we added as the default

`kops validate cluster`

it will show whether cluster is ready or not

`ssh -i .ssh/id_rsa admin@<ip address of master or node>`

from here we can connect to master or node through ssh connection

`cat .kube/config`

it will show all the information of cluster

Creating a Basic pod

`kubectl run <pod name> --image=index.docker.io/<repository name>/<image name>:v1`

it will create the pod by giving repository name and image from docker hub

here index.docker.io = docker hub

here run cmd will process deployment and we can add --replicas=3(it will create replicas)

```
kubectl get pods
```

here it will show all pods

```
kubectl describe pod <podname>
```

it will give detail description of a pod and what is happening in background

```
kubectl get pod <podname> -o yaml
```

it will show detail description of a pod with ip address, restarts, status....etc. in yaml format

```
Kubelet expose pod <podname> --port=8000 --target-port=80 --type=Nodeport
```

it will expose the node to specific port

```
kubectl get svc
```

it will show the node on which port it is exposed

```
kubectl edit svc <Node Name>
```

it will edit service of specific node

```
kubectl delete svc <Node Name>
```

it will delete service for specific port

```
kubectl get all
```

it will the list of pod

```
kubectl exec -it <podname> -- ls -al
```

it will show all directories in pod

```
kubectl exec -it <podname> -- bash
```

we can enter into the specific pod

```
kubectl delete pod <podname1> <podname2>
```

it will delete the specific pods

```
kubectl create -f pod.yaml
```

it will create the pod with the specific yaml file content

```
kubectl apply -f pod.yaml
```

if any changes to be done on the created pod we need to give apply in the cmd

```
kubectl ns <name>
```

it will create a namespace for the given

```
kubectl get pods -l env=dev
```

here in yaml we can mention labels, e.g.; env=dev,prod...etc. so while executing the cmd we can mention the label name(env=dev) it will sort the specific pods with that label

kubectl get pods --show-labels

it will all labels of every pod

kubectl label node <nodename> app-

here it will remove the label of the pod, If you're running a replicas then a new replica will be created

kubectl get rc

it will show the replication controller pods

kubectl delete rc <ReplicationControllerName>

it will delete the all pods within the Replication controller name

kubectl expose rc <ReplicationControllerName> --port=8000 --target-port=80 --type=NodePort

it will expose the replication controller to specific port

kubectl expose deploy <DeployName> --port=8000 --target-port=80 --type=NodePort

it will expose the Deploy to specific port

**kubectl rolling-update <ReplicationControllerName> <New Name> --
image=index.docker.io/sreeharshav/testcontainer:v1**

here rolling-update is for changing the version of application without any downtime... it will upgrade one by one pod

kubectl rolling-update <ReplicationControllerName> <New Name> --rollback

here it will take back to old version

Deployment Advantages:

1. it is uses replicaset and replicaset automatically perform the rollingupdates
2. Rollback is easy as we can record the deployments
3. we can use liveness and readiness probes to improve application availability.
4. We can Pause & resume the deployment which useful canary update.

kubectl set image deployment <Deployment Name> nginx=<repository name>/<image name>:v1

this will replace our deployed image to the new image, but it was not best practice we need to change image in yaml

kubectl describe deploy <Deployment Name>

it will all details of deployment... eg: image used for deployment

kubectl scale deployment <deployment Name> --replicas=5

it will increase the replicas

kubectl rollout history deployments <deployment Name> --revision=2

here we should give which revision.... so it will show the changes done in the specific revision

Rolling Update

```
---
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  minReadySeconds: 20
  replicas: 3
  selector:
    matchLabels:
      maxsurge: 1
      maxUnavailable: 0
    type: RollingUpdate
  template:
    metadata:
      labels:
        app: nginx
    specs:
      container:
        - name: nginx
          image: sreeharshav/testcontainer:v1
          ports:
            - containerPort: 80
---
kind: Service
apiVersion: v1
metadata:
  name: myservice
spec:
  selector:
    app: nginx
  type: NodePort
  ports:
    - name: name-of-the-port
      port: 8000
      targetPort: 80
```

Probes: Health checks

1.Rediness Probe => check the POD is ready to service the requests then only provide traffic

2.Liveness Probe => Check the POD is responding, if not restart it

here KUBELET will perform the probes.

```
specs:
  container:
    - name: nginx
      image: nagasiva/testcontainer:v1
      ports:
        - containerPort: 80
      readinessProbe:
        initialDelaySeconds: 30
        PeriodSeconds: 5
        timeoutSeconds: 10
        successThreshold: 1
        failureThreshold: 3
        httpGet:
          path: /
          Port: 80
      livenessProbe:
        initialDelaySeconds: 60
        PeriodSeconds: 5
        timeoutseconds: 10
        successThreshold: 1
        failureThreshold: 1
        httpGet:
          path: /
          Port: 80
```

kubect! get events

it will show all the updates done on the node or pod and it will show where readiness and liveness failed

Types of services

here we have three services:

- i. Node port
- ii. Cluster Ip
- iii. None(Headless Service)
- iv. Load Balancer
- v. ingress

- I. **Nodeport:** It is a service which exposes the node directly, it can be accessed from internet (It is to communicate with Single node)
- II. **Cluster IP:** It will enable the communication between node to node internally within the cluster, it will be not exposed to internet, by default if we create service cluster ip will be taken (from frontend application to backend)

- III. **None(Headless Service):** it is service which it can given to one backend pod like databases, it is used when you don't need load balancer(it is not exposed to internet)
- IV. **Load Balancer:** A Load Balancer service is the standard way to expose a service to the internet, but it can be used for only one service
- V. **Ingress:** it is similar to load Balancer but it is used for Different Services

NodePort:

kubectl expose pod <pod name> --port=8000 --target-port=80 --type=nodeport

```
apiVersion: v1
kind: Service
metadata:
  name: mysql
  labels:
    name: mysql
spec:
  type: NodePort
  ports:
    - port: 3036
      nodePort: 30036
      name: http
  selector:
    name: mysql
```

Cluster IP:

here no need of giving "type" by default it will take as Cluster IP

kubectl expose deploy <deployment name> --port=8000 --target-port=80

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
spec:
  selector:
    app: webserver-nginx
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080
```

None(Headless Service):

kubectl expose <deployment name> --port=8000 --target-port=80 --cluster-ip=None

```
apiVersion: v1
kind: Service
```

```

metadata:
  name: headless-service
spec:
  clusterIP: None # <-- Don't forget!!
  selector:
    app: api
  ports:
    - protocol: TCP
      port: 80
      targetPort: 3000

```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/headless-service	ClusterIP	None	<none>	80/TCP	5s
service/kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	45h
service/normal-service	ClusterIP	10.109.192.226	<none>	80/TCP	5s

Load Balancer:

kubectl expose deploy <deployment name> --port=8000 --target-port=80 --type=LoadBalancer

```

apiVersion: v1
kind: Service
metadata:
  name: my-cip-service
spec:
  type: ClusterIP
  selector:
    app: metrics
    department: sales
  ports:
    - protocol: TCP
      port: 80
      targetPort: 8080

```

Ingress:

we need to create NGINX Ingress Controller

```
$ kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v0.44.0/deploy/static/provider/cloud/deploy.yaml
```

```

apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: simple-fanout-example
spec:
  rules:
    - host: foo.bar.com

```

```

http:
  paths:
  - path: /foo
    pathType: Prefix
    backend:
      service:
        name: service1
        port:
          number: 4200

```

`kubectl apply -f <.yaml>`

it will apply the service

`kubectl get ing`

it will show all ingress services

Name Spaces:

- Allowing teams or projects to exist in their own virtual clusters without fear of impacting each other's work.
- Enhancing role-based access controls (RBAC) by limiting users and processes to certain namespaces.
- Enabling the dividing of a cluster's resources between multiple teams and users via resource quotas.
- Providing an easy method of separating development, testing, and deployment of containerized applications enabling the entire lifecycle to take place on the same cluster.

First need to download kubens from internet

<https://github.com/ahmetb/kubectx/releases>

take the latest version

and unzip the file

`mv kubens /usr/local/bin/`

move the folder to specific directory

`kubectl get ns`

It will show all namespaces list

```

root@ip-192-168-1-100:~# ku get ns
NAME                STATUS    AGE
default             Active    26m
kube-node-lease     Active    26m
kube-public         Active    26m
kube-system         Active    26m

```

if you didn't create any namespace it will go to default namespace

```
kubectl create ns <namespace name>
```

it will create a namespace with the given name

```
kubectl delete ns <namespace name>
```

it will delete the namespace

```
kubens <Name of Namespace>
```

this command will convert given specific namespace to Default

```
kubectl apply -f pod.yaml --namespace=test
```

It will create pod with the specific test namespace

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
  namespace: test
  labels:
    name: mypod
spec:
  containers:
  - name: mypod
    image: nginx
```

```
kubectl describe ns <name>
```

It will show all details of the namespace like-status, resource quota, limit range resource.

ResourceQuota:

It is for limiting the total resource consumption of a namespace

limit range:

It is for managing constraints at a pod and container level within the project

ResourceQuota:

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: compute-resources
spec:
  hard:
    pods: "4"
    requests.cpu: "1"
    requests.memory: 1Gi
    limits.cpu: "2"
    limits.memory: 2Gi
```

`kubectl apply -f <yaml file of Resourcequotas> --namespace=<Name of namespace>`

it will create the Resourcequota for the specific namespace

`kubectl get resourcequotas`

it will show all resourcequotas created

`kubectl delete resourcequotas <name>`

it will delete the specific quota

limit range:

```
apiVersion: "v1"
kind: "LimitRange"
metadata:
  name: "core-resource-limits"
spec:
  limits:
    - type: "Pod"
      max:
        cpu: "2"
        memory: "1Gi"
      min:
        cpu: "200m"
        memory: "6Mi"
    - type: "Container"
      max:
        cpu: "2"
        memory: "1Gi"
      min:
        cpu: "100m"
        memory: "4Mi"
      default:
        cpu: "300m"
        memory: "200Mi"
      defaultRequest:
        cpu: "200m"
        memory: "100Mi"
      maxLimitRequestRatio:
        cpu: "10"
```